

# Overwriting the Exception Handling Cache Pointer - Dwarf Oriented Programming

James Oakley

Electron100 \*noSPAM\* gmail.com

Rodrigo Rubira Branco (@BSDaemon)

rodrigo \*noSPAM\* kernelhacking.com

Sergey Bratus (@SergeyBratus)

Sergey \*noSPAM\* cs.dartmouth.edu

# Credits

- This presentation combines ideas, research, discussions from the following personnel:
  - Sergey Bratus (“Insecurity Theory”, Exploiting the Hard-working Dwarf)
  - Meredith Patterson (Langsec)
  - R.I.P. Len Sassaman (Langsec)
  - James Oakley (Exploiting the Hard-working Dwarf -> everything related to that, including Katana)
  - Rodrigo Rubira Branco (Exploiting the Hard-working Dwarf -> exploitation, implementation details, research organization)

# Motivation

- Software exploitation is **not generic** anymore
- There are different exploitation primitives in different contexts
- A modern exploitation technique shows how to take advantage of those primitives
- Targets as written are capable of many more computations than intended. **Exploitation is proof of that**

# A Bit of Theory

- *Trustworthiness* of a computer system is **what the system can and cannot compute**
  - Can the system decide if an input is invalid/unexpected/malicious & reject it?
  - Will program perform only *expected* computations, or *malicious* ones too?
- ***Exploitation is setting up, instantiating, and programming a weird machine***
  - A part of the target is overwhelmed by crafted input and enters an **unexpected** but **manipulable** state

# Software Exploitation

- A part of the target is overwhelmed by crafted input and enters an **unexpected** but **manipulable** state
- **Primitives** are exposed
  - Memory corruption, unexpected control flows, ...
  - From a formal model point of view, primitives supply extra state and/or transitions
- A “weird machine” is unleashed
  - A **more powerful**, programmable **execution environment** than intended or expected

# Weird Machine is Born!



# Exploiting Additional Computations

- Finally we are in our talk line...
- There are many computations inside a program that can be used to subvert the code execution (and some of them has nothing to do with the original code itself)
- ROP is not new, exploits are using it since 2000 (maybe even before)

# \*nix Exception Handling

- Binaries compiled with GCC and that support exception handling have Dwarf bytecode:
  - Describe the stack frame layout
  - Interpreted to unwind the stack after an exception occurs
- The process image includes the Dwarf interpreter (part of the GNU C++ runtime)
- Bytecode can be written to force the interpreter to perform any computation (Turing-Complete), including, but not limited to, setup a library/system call modifying registers such as stack and base pointers -> See James and Sergey previous work on Dwarf Trojans



# James Oakley and Sergey Bratus

- Proved that Dwarf can replace code creating a Trojan completely using Dwarf bytecode
- Proved that Dwarf is a complete development environment:
  - Can read memory
  - Can compute with values from memory/registers
  - Can influence the flow of execution of a process

# ELF

ELF Header
Program Headers
<b>.init</b>
<b>.plt</b>
<b>.text</b>
<b>.fini</b>
<b>.eh_frame_hdr</b>
<b>.eh_frame</b>
<b>.gcc_except_table</b>
<b>.dynamic</b>
<b>.got</b>
<b>.data</b>
<b>.symtab</b>
<b>.strtab</b>
Section Headers

The executable has this format either on disk or in memory.

# Dwarf

- Developed as a debugging format to replace STABS
- Standard: <http://dwarfstd.org>
- Provide information such as code line, variable types, backtraces, others
- ELF Sections: .debug\_info, .debug\_line, .debug\_frame are defined in the standard
- .debug\_frame defines how to unwind the stack (how to restore each entry in the previous call frame)

# Linux Exception Handling

- GCC, the Linux Standards Base and the ABI x86\_64 adopted a very similar format used in the `.debug_frame` to describe the stack unwind during an exception: `.eh_frame`
- It is not exactly the same as dwarf
- It adds pointer encoding and language-specific data
- As usual, the documentation is not perfect:
  - Partially discussed in the Linux Standards Base
  - Partially defined in the ABI
  - Partially implemented in GCC

# .eh\_frame

- Theoretically it is a table, where for each address in the .text it is describe how to restore the registers to the previous call frame

EIP	CFA	EBP	EBX	EAX	RET
0xf000f000	rsp+16	*(cfa-16)			*(cfa-8)
0xf000f001	rsp+16	*(cfa-16)			*(cfa-8)
0xf000f002	rbp+16	*(cfa-16)		eax=edi	*(cfa-8)
...	...	...	...	...	...
0xf000f00a	rbp+16	*(cfa-16)	*(cfa-24)	eax=edi	*(cfa-8)

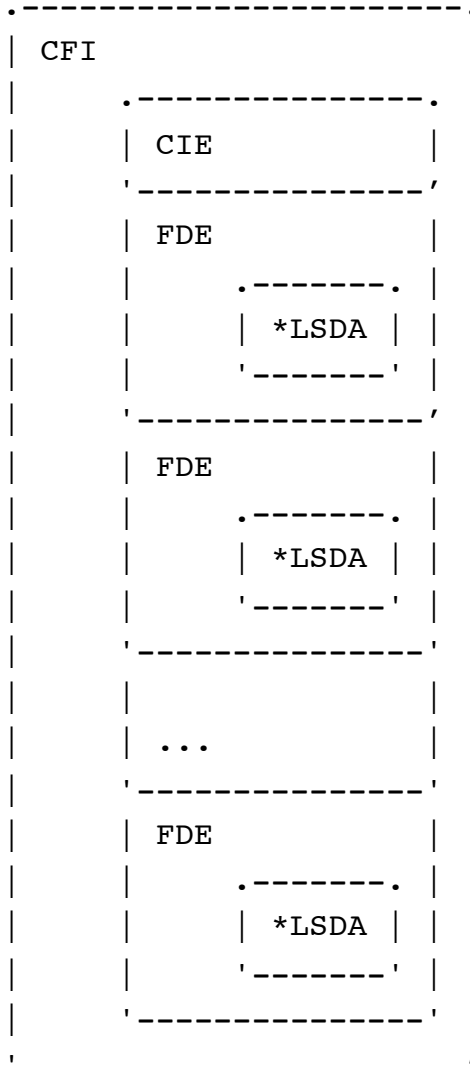
- CFA (Canonical Frame Address) – Address relative to the call frame
- Each line defines how each part of the code can return to the previous frame

# Size Limitations

- Obviously, keep such a table would use more space than the code itself
- That's why the adoption of bytecode: The table is 'compressed', providing everything required to create it when needed
- Portions of the table are created as needed (on-demand)

# .eh\_frame

.eh\_frame section



CFI = Call Frame Information

CIE = Common Information Entry

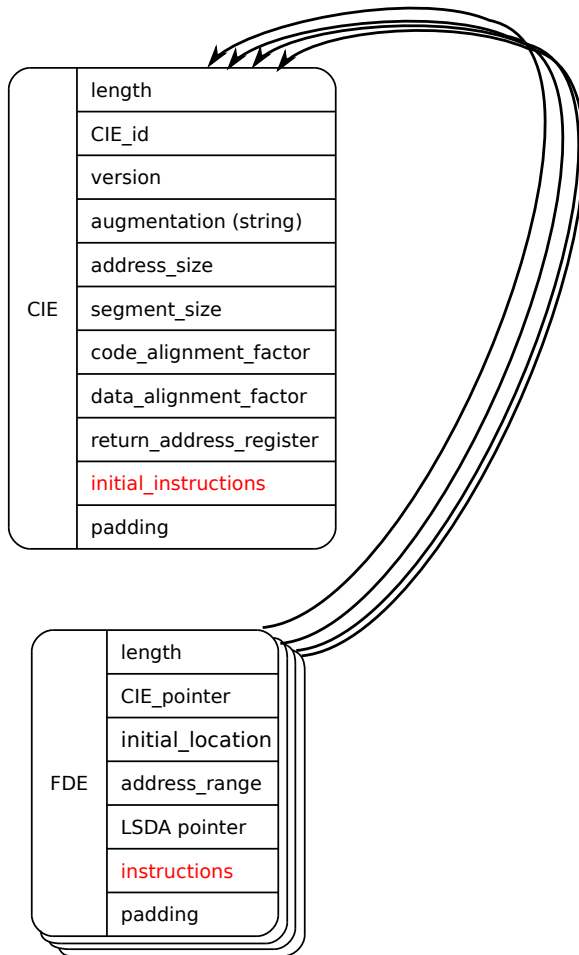
FDE = Frame Description Entry

LSDA = Language Specific Data Area

"The .eh\_frame section shall contain 1 or more Call Frame Information (CFI) records. The number of records present shall be determined by size of the section as contained in the section header. Each CFI record contains a Common Information Entry (CIE) record followed by 1 or more Frame Description Entry (FDE) records. Both CIEs and FDEs shall be aligned to an addressing unit sized boundary"

***Linux Standard Base Specification 1.3***

# FDE x CIE




FDE (Frame Description Entry) exists for each logical Instruction block

CIE (Common Information Entry) holds common Information between FDEs

INSTRUCTIONS in FDE hold the DWARF bytecode



# FDE x CIE



## **initial\_location/address\_range:**

Defines for which instructions this FDE applies

## **augmentation:**

Language-specific information

## **return\_address\_register:**

Entry in a virtual table that defines the .text location to return to (eip)

## **instructions:**

Table rules. Dwarf has a language to describe the table.

# Dwarf Instructions

- Work as an assembly language (unexpected computations)
- Turing-Complete Stack-Based Machine
- Can access memory and register values
- Have some limitations:
  - Cannot write to register/memory (but we can force out-of-order code execution and obtain writes)
  - Cannot call native code
  - Cannot write to registers that are not callee-saved in the ABI (we can write to callee-saved register though)
  - GCC limits the stack to 64 words

# Dwarf Programming

- **DW\_CFA\_set\_loc N**  
Next instructions apply to the first N bytes of the function
- **DW\_CFA\_def\_cfa R OFF**  
CFA is calculated starting from register R and offset OFF
- **DW\_CFA\_offset R OFF**  
Register R is restored from the value in CFA OFF
- **DW\_CFA\_register R1 R2**  
Register R1 is restored with the contents of R2

# And the table is back...

- Each architecture register receives a DWARF equivalent (the mapping is architecture specific)
- Dwarf Instructions define rules for a column or advances to the next line (program location)
- In a FDE, lines heritage from instruction lines above them

EIP	CFA	EBP	EBX	EAX	RET
0xf000f000	rsp+16	*(cfa-16)			*(cfa-8)
0xf000f001	rsp+16	*(cfa-16)			*(cfa-8)
0xf000f002	rbp+16	*(cfa-16)		eax=edi	*(cfa-8)
...	...	...	...	...	...
0xf000f00a	rbp+16	*(cfa-16)	*(cfa-24)	eax=edi	*(cfa-8)

# Dwarf Expressions

- To not anticipate all unwinding mechanisms of a system, the standard defines flexibility:
  - DW\_CFA\_expression R EXPRESSION  
R receives the value from the EXPRESSION result
  - DW\_CFA\_val\_expression R EXPRESSION  
R restored to result of EXPRESSION
- Expressions have their own instructions:
  - Constant Values: DW\_OP\_constu, DW\_OP\_const8s, etc
  - Arithmetic: DW\_OP\_plus, DW\_OP\_mul, DW\_OP\_and, DW\_OP\_xor, etc
  - Memory read: DW\_OP\_deref
  - Register read: DW\_OP\_bregx
  - Flow Control: DW\_OP\_le, DW\_OP\_skip, DW\_OP\_bra, etc

# Katana

## **Emit a dwarfscript**

➤ \$e=load "demo"

Loaded ELF "demo"

➤ dwarfscript emit ".eh\_frame" \$e "demo.dws"

Wrote dwarfscript to demo.dws

## **Dwarfscript assembler**

➤ \$ehframe=dwarfscript compile "demo.dws"

➤ replace section \$e ".eh\_frame" \$ehframe[0]

Replaced section ".eh\_frame"

➤ save \$e "demo\_rebuilt"

Saved ELF object to "demo\_rebuilt"

➤ !chmod +x demo\_rebuilt

# So what?

- With Katana you can see and modify unwind tables in an easy way
  - Control the unwinding flow (how the call stack is handled)
  - Avoid an exception handler to execute another one
  - Redirect exceptions
  - Find/solve symbols
  - Calculate relocations

# Example

- If function foo is responsible for an exception
  - Change flow to function bar
  - Thru static analysis, we see that bar is at 0x600DF00D
  - In the FDE, we change:  
    DW\_CFA\_offset r16 1
  - To:  
    DW\_CFA\_val\_expression r16  
    begin EXPRESSION  
    DW\_OP\_constu 0x600DF00D  
    dnd EXPRESSION



# `.gcc_except_table`

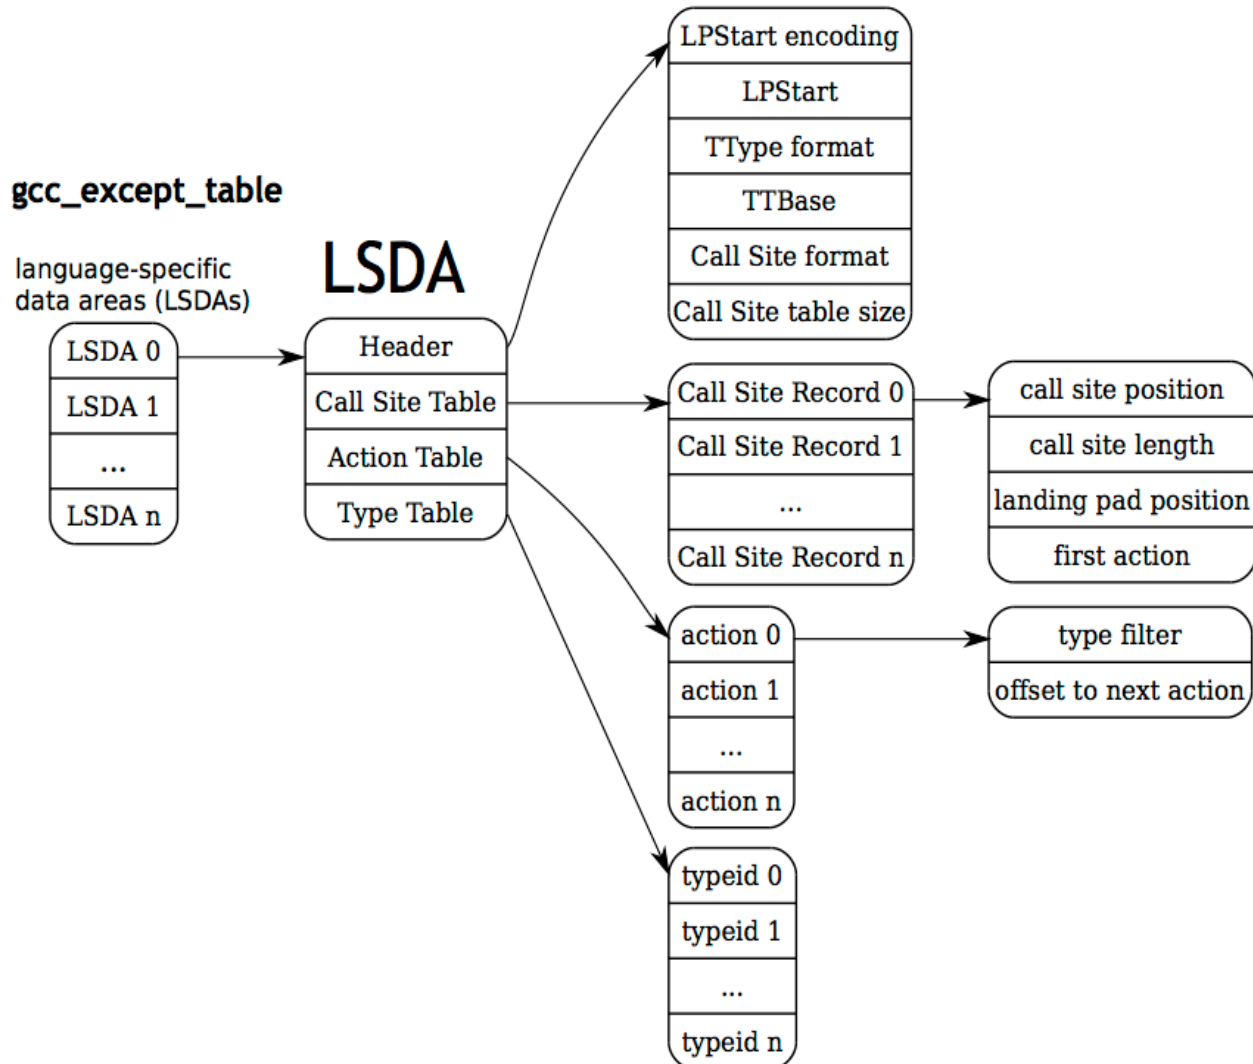
- So far, redirected only to 'catch' blocks
- The `.gcc_except_table` hold language-specific data (where the exception handlers are)
  - Interpreted by the personality routines
  - We can stop an exception at any time
  - Unlike the `.eh_frame`, do not have standards
  - There is no documentation, so let's see the code ;)

# Assembly

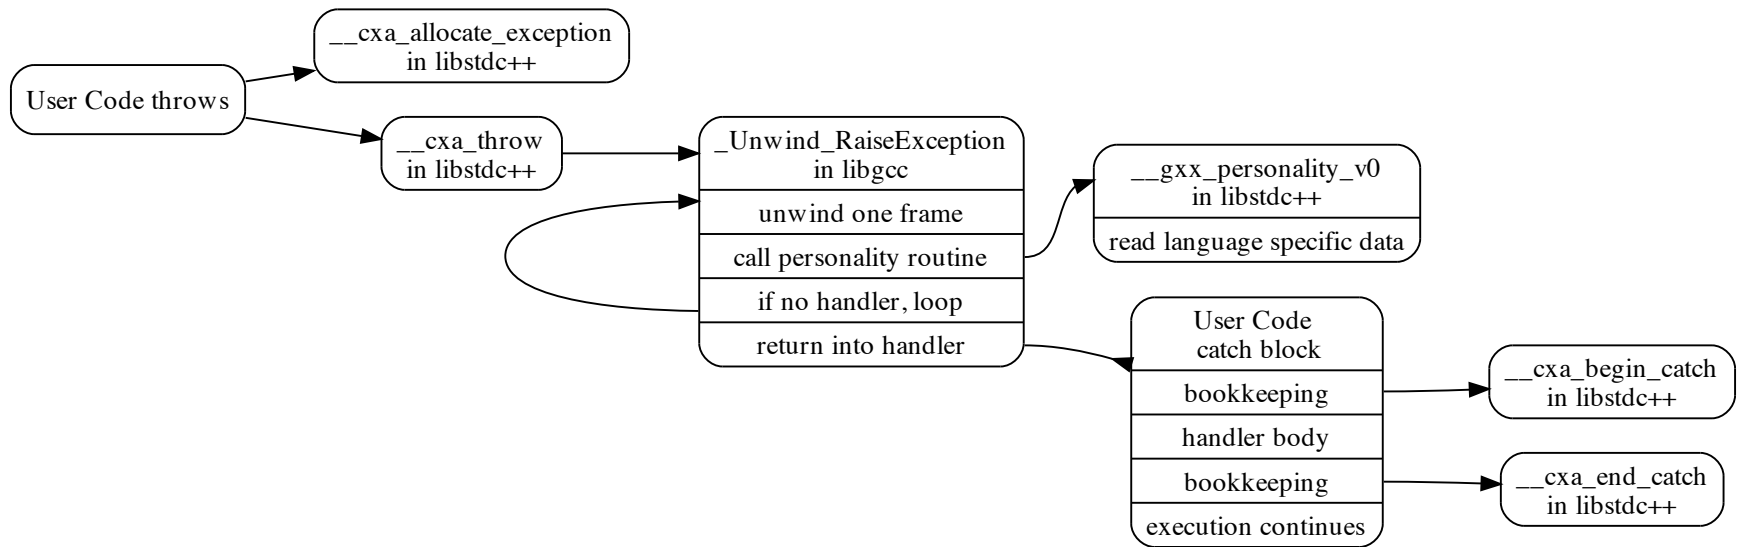
- While compiling a program using GCC, do:  
--save-temps -fverbose-asm -dA

```
.section .gcc_except_table,"a",@progbits
.align 4
.LLSDA963:
    .byte 0xff # @LPStart format (omit)
    .byte 0x3 # @TType format (udata4)
    .uleb128.LLSDATT963-.LLSDATTD963 # @TType base offset
.LLSDATTD963:
    .byte 0x1 # call-site format (uleb128)
    .uleb128 .LLSDACSE963-.LLSDACSB963 # Call-site table length
.LLSDACSB963:
    .uleb128 .LEHB0-.LFB963 # region 0 start .uleb128 .LEHE0-.LEHB0 #
length .uleb128 .L6-.LFB963 # landing pad .uleb128 0x1 # action
    .uleb128 .LEHB1-.LFB963 # region 1 start .uleb128 .LEHE1-.LEHB1 #
length .uleb128 0x0 # landing pad
    .uleb128 0x0 # action
    .uleb128 .LEHB2-.LFB963 # region 2 start .uleb128 .LEHE2-.LEHB2 #
length .uleb128 .L7-.LFB963 # landing pad .uleb128 0x0 # action
.LLSDACSE963:
    .byte 0x1 # Action record table .byte 0x0
    .align 4
    .long _ZTIi
```

# Layout



# Exception Handling Flow



# Exceptions are not asynchronous

- Functions that call `throw()` just call:
  - `__cxa_allocate_exception()` -> To allocate space using `malloc` (or buffers in the `.bss` if `malloc` fails – `gcc-xxx/libstd++v3/libsupc++/eh_alloc.:84`)
  - And then `__cxa_throw()` -> That will go thru the frames until a handler for the exception is found

# Proving (assembly)

Dump of assembler code for function main:

...

```
<+9>:  mov    $0x4,%edi          # std::size_t thrown_size
```

# Allocates a new "\_\_cxa\_refcounted\_exception" followed by 4 bytes; we

# do a "throw(1)", 1 being an "int" occupies 4 bytes.

```
<+14>:  callq  0x400930 <__cxa_allocate_exception@plt>
```

...

```
<+25>:  mov    $0x0,%edx           # void (*dest) (void *)
```

```
<+30>:  mov    $0x6013c0,%esi      # std::type_info *tinfo
```

```
<+35>:  mov    %rax,%rdi           # void *obj
```

```
<+38>:  callq  0x400940 <__cxa_throw@plt>
```

# \_\_cxa\_allocate\_exception()

- Returns a pointer to a
  - struct \_\_cxa\_refcounted\_exception, which holds a reference to an object \_\_cxa\_exception
- \_\_cxa\_throw() is then executed to:
  - Initialize the current context (register values)
  - Iterate in the stack until it finds the exception handler

# What We've Shown Before

- Ret-into-libc
- Used the dynamic-linker already in Dwarf to find `execvpe`
- Used Dwarf to prepare the stack
- In less than 200 bytes and less than 20 words in the stack (showing that a 64-stack word limitation is not an obstacle)
- Started in an offset of `execvpe` where they can control the Dwarf registers (and not in the function beginning)



# What else can be done?

- Old GCC had both, the `.eh_frame` and the `.gcc_except_table` as +W
- Well...
  - Libgcc/libstdc++ need to find those areas in memory, right?
  - The program header, `GNU_EH_FRAME` contains the `.eh_frame` location (`dl_iterate_phdr` is the function that finds it)
  - Libgcc caches the value!

# Fake EH

- If we can overwrite the cached value, we are able to control the exceptions and leverage everything already explained
- Libgcc does not export symbols, so we need to find an heuristic/reverse to find what to overwrite

# Caching

- The pointer caching is done in: unwind-dw2-fde-glibc.c:

```
#define FRAME_HDR_CACHE_SIZE 8
```

```
...
```

```
static struct frame_hdr_cache_element
```

```
{
```

```
    _Unwind_Ptr pc_low;
```

```
    _Unwind_Ptr pc_high;
```

```
    _Unwind_Ptr load_base;
```

```
    const ElfW(Phdr) *p_eh_frame_hdr;
```

```
    const ElfW(Phdr) *p_dynamic;
```

```
    struct frame_hdr_cache_element *link;
```

```
} frame_hdr_cache[FRAME_HDR_CACHE_SIZE];
```

# Caching

- 8 cache entries for the frame header
  - Uses a Least Used Replacement Algorithm (`_Unwind_IteratePhdr_Callback()`)
  - Most recently used is the head of the list
- In the test environment, the `frame_hdr_cache` was at 0x6e0 bytes from the offset of the writable data segment of `libgcc`
- This is the aforementioned array, with 48 bytes in size
- The executable itself is the 3<sup>rd</sup> element of the array (the first two are the `libgcc` and `libstdc++`)
- The offset for the writable data segment of `libgcc` can be found in this way (based in what we know):
  - $0x6e0 + 48 * 2 = 0x740$
- The entry `p_eh_frame_hdr` that we want to overwrite is at 24 bytes of this structure.

# Exploiting

- To simplify the exploitation, it is interesting to align the structures in known offsets/  
controlled offsets:
  - .eh\_frame in the example aligned to start exactly at 0x50 bytes from the start of the .eh\_frame\_hdr
  - .gcc\_except\_table aligned to start exactly at 0x200 bytes from the start of the .eh\_frame

# Memleak

- We need the value of EBP, so we going to use a memleak. It can be achieved in different ways, depending on the target program (e.g.: overwriting parameters to printf-like functions, or if the vulnerability is a format string, which is our sample case)
- To calculate the EBP\_PREVIOUS we use %llx (format string), so we use 4 bytes of space in the buffer and advance the stack pointer in 8 bytes (so the premise for exploiting the sample program is to manage to leak the EBP):

• • •

```
#to get the value of ebp_previous
```

[illegible]

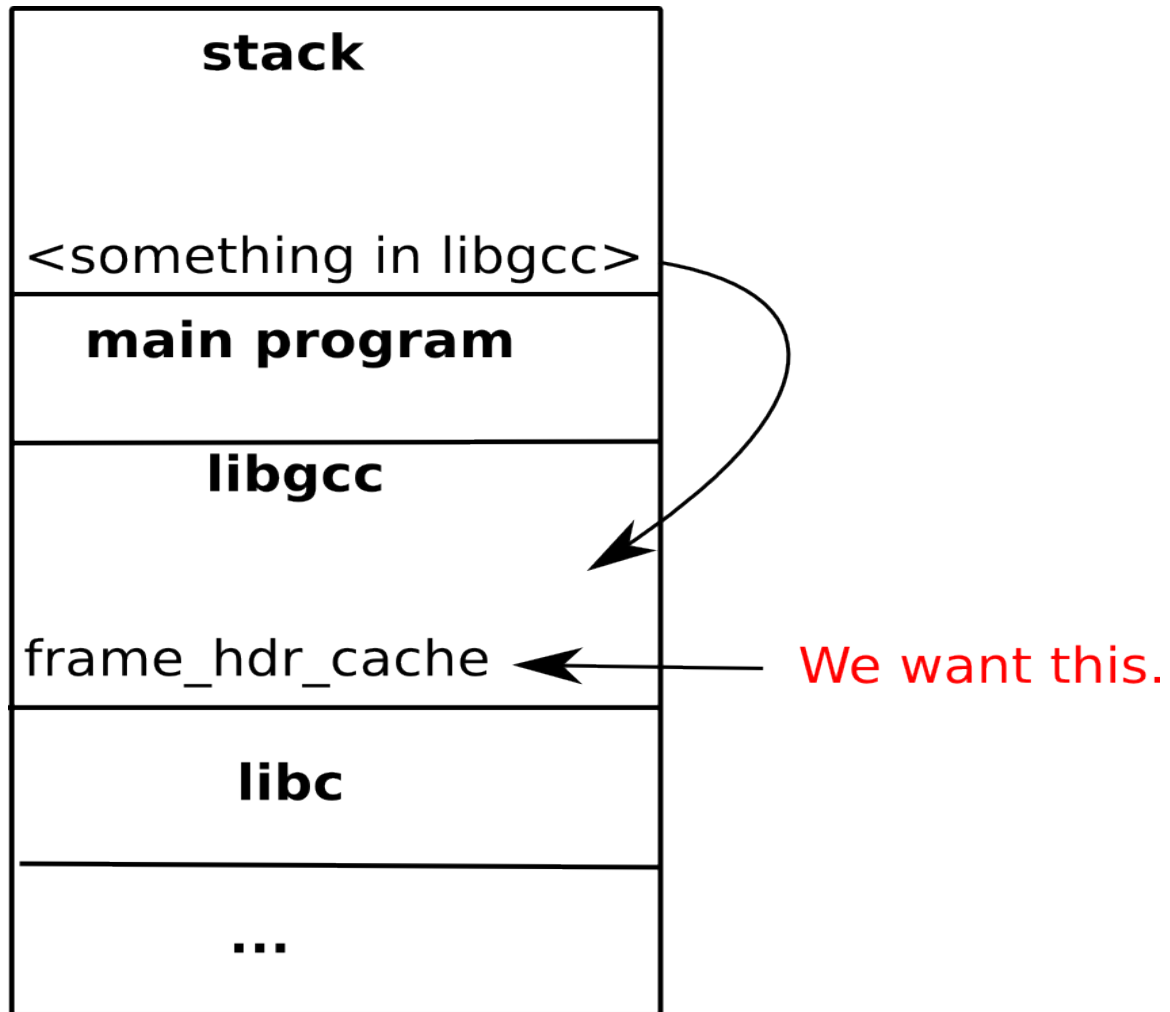
```
proc.sendline(instr)
```

```
proc.expect("unknown command: [0-9a-f]* ([0-9a-f]*)*.*)")
```

```
ebp previous=int(proc.match.group(1),16)
```

```
info("\nfound ebp_previous = 0x%x" % ebp_previous)
```

# Memleak



# Heuristics

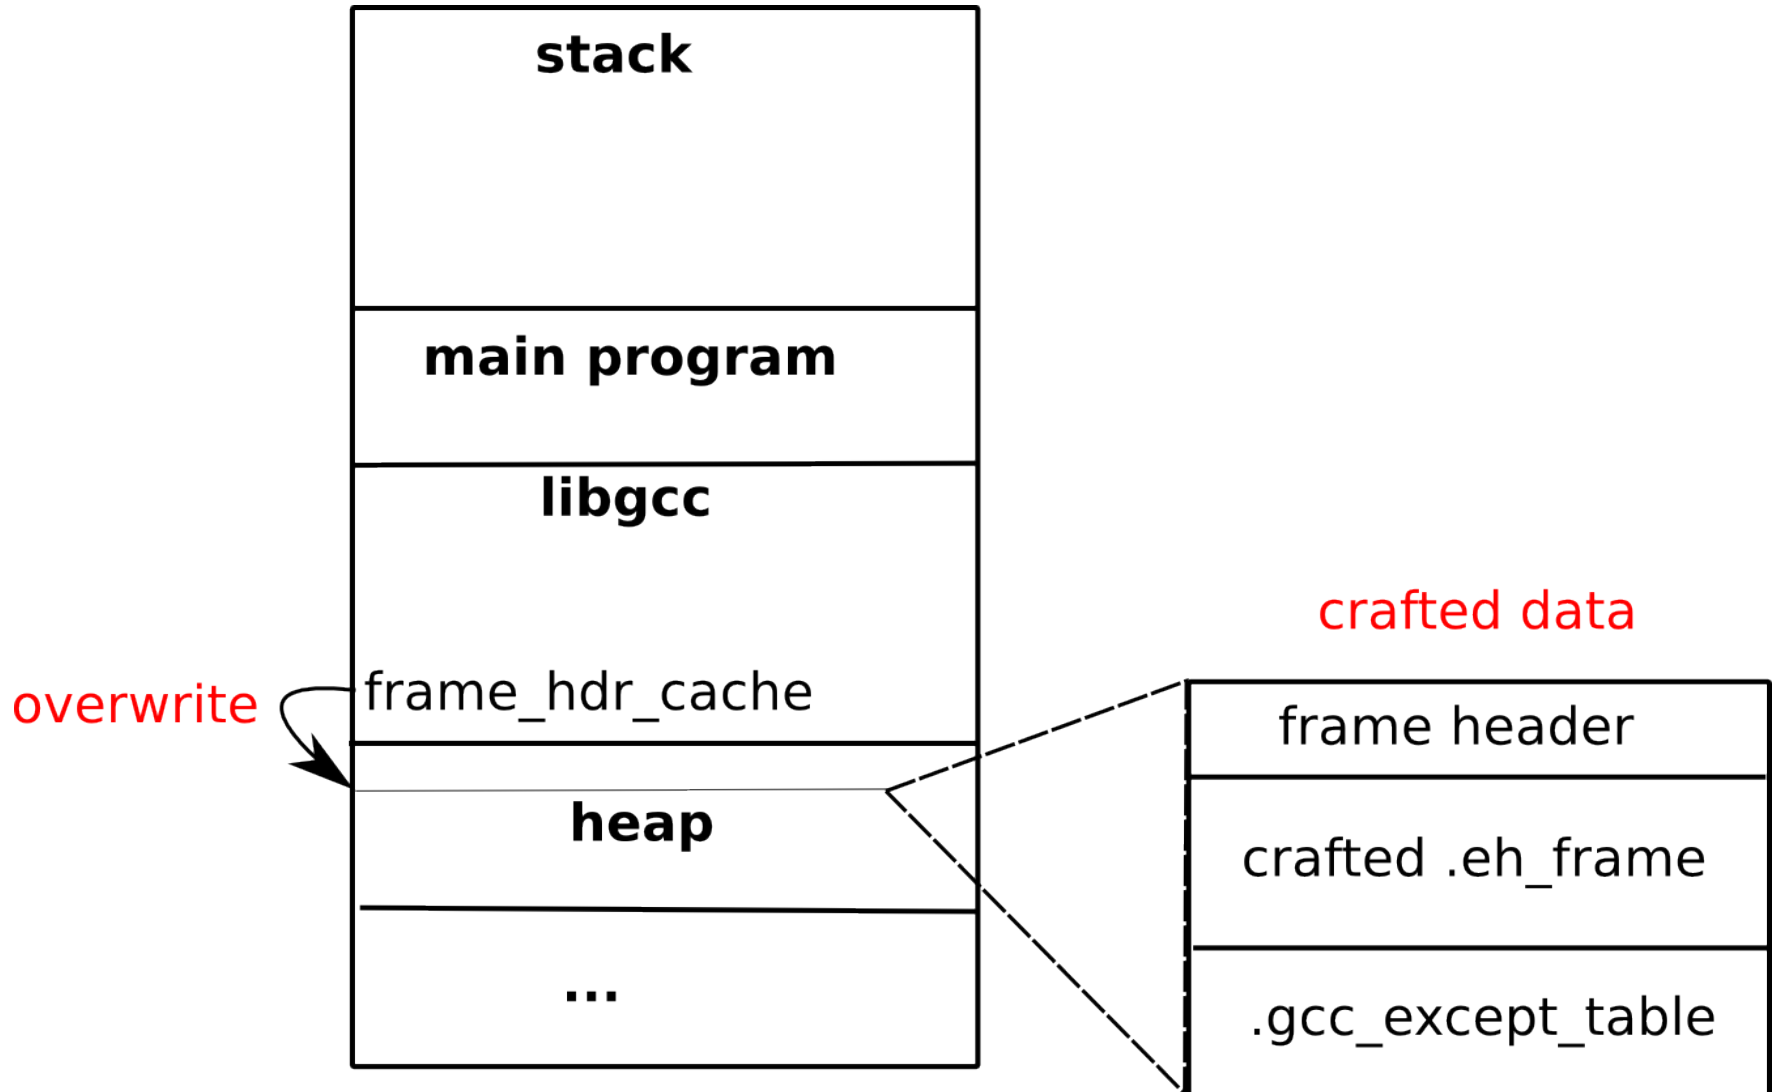
- We know the size of the previous frame (disassembling), so we are capable of calculate the EBP of our frame:
  - `ebp=ebp_previous-PREV_FRAME_SIZE`
- With our frame address, we can calculate the address of libgcc, since we know the offsets:
  - `libgcc_reveal_location=ebp-LIBGCC_REVEAL_EBP_OFFSET;`



# More Heuristics

- The value that reveals the .text location of the libgcc is at 0xffffc798 (discovered in the previous slide), and it is 0x679 above ESP and 0x750 above EBP
- The libgcc base is calculated using the previously revealed address and masking the 3 low nibbles. We also use a fixed value to adjust the result (found thru disassembly):
  - `libgcc_base=(libgcc_revealed & 0xFFFFF000) - LIBGCC_REVEAL_ADJUST`
- The separation between .text and .data segments in libgcc is 0x19000 (x86):
  - `libgcc_data_base=libgcc_base+LIBGCC_DATA_OFFSET`

# Overwriting the Frame Header Cache



# Finalizing

- Finally, we find the `frame_hdr_cache` and the respective `p_eh_frame_hdr` from the `libgcc_data_base`, as previously described:
  - `frame_hdr_cache=libgcc_data_base+CACHE_LIBGCC_OFFSET`
  - `p_eh_frame_hdr=frame_hdr_cache+CACHE_ENTRY_SIZE*PREVIOUS_CACHE_ENTRIES+OFFSET_IN_CACHE_ENTRY`

# In the demo case

- The Dwarf payload is injected in the dictionary been readed by the target program (instead of using a shellcode). We find the pointers, overwrite the caching target address and the desired catch block is executed.
- With all the values, we redirect the execution:
  - Function doWork throws an exception
  - We redirect it to I\_am\_never\_called

# Other possibilities

- If you have a Write N you can overwrite the .eh\_frame entirely (if it is +W, what is not normal in new systems)
- You can overwrite the .eh\_frame using a shellcode
- You can use a staggered ret-into-lib to remap the .eh\_frame as +W and then overwrite it

# THE END! Really is !?

<http://katana.nongnu.org>

<https://github.com/rrbranco/defcon2012>

James Oakley (Electron)

Electron100 \*noSPAM\* gmail.com

Rodrigo Rubira Branco (@BSDaemon)

rodrigo \*noSPAM\* kernelhacking.com

Sergey Bratus (Sbratus)

Sergey \*noSPAM\* cs.dartmouth.edu